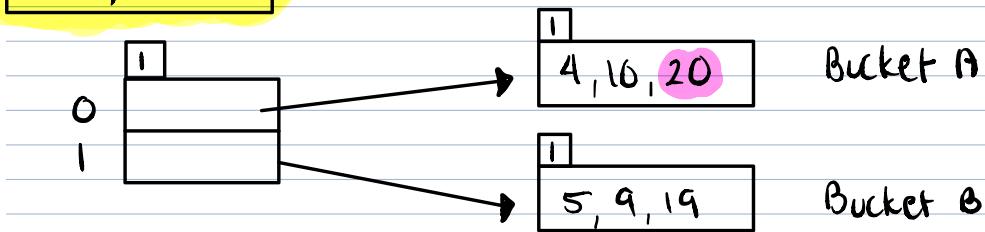


Jerry Perez

Problem 1

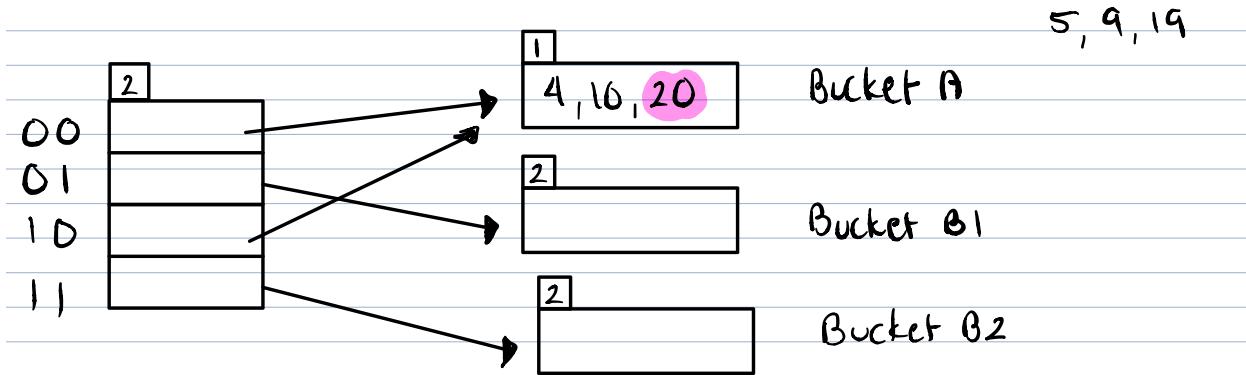


Binary conversions

32 16 8 4 2 1

insert 20 0 1 0 1 0 0, Bucket A

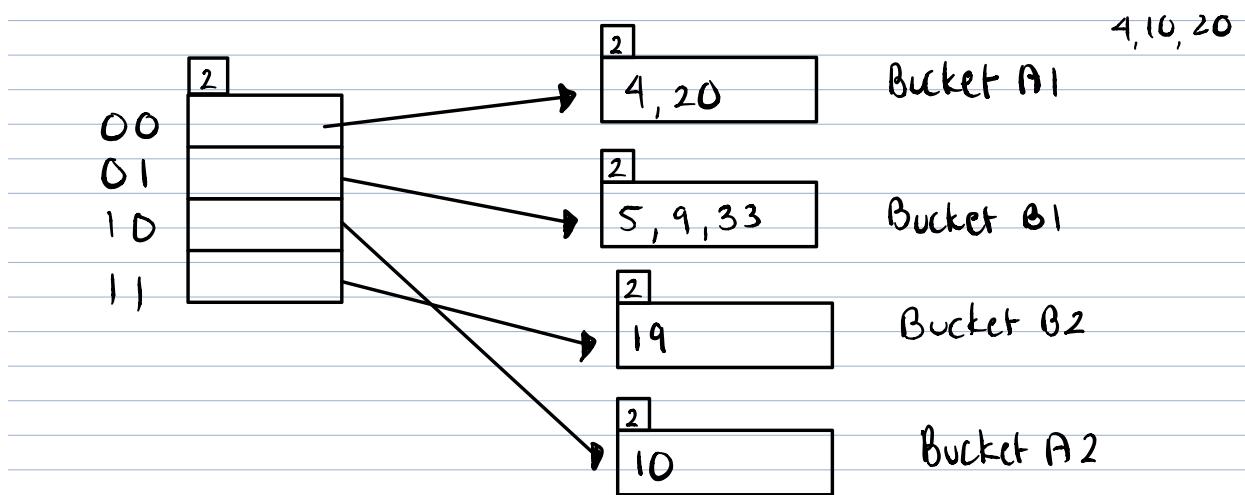
insert 33 1 0 0 0 0 1, Bucket B-split



Binary conversions

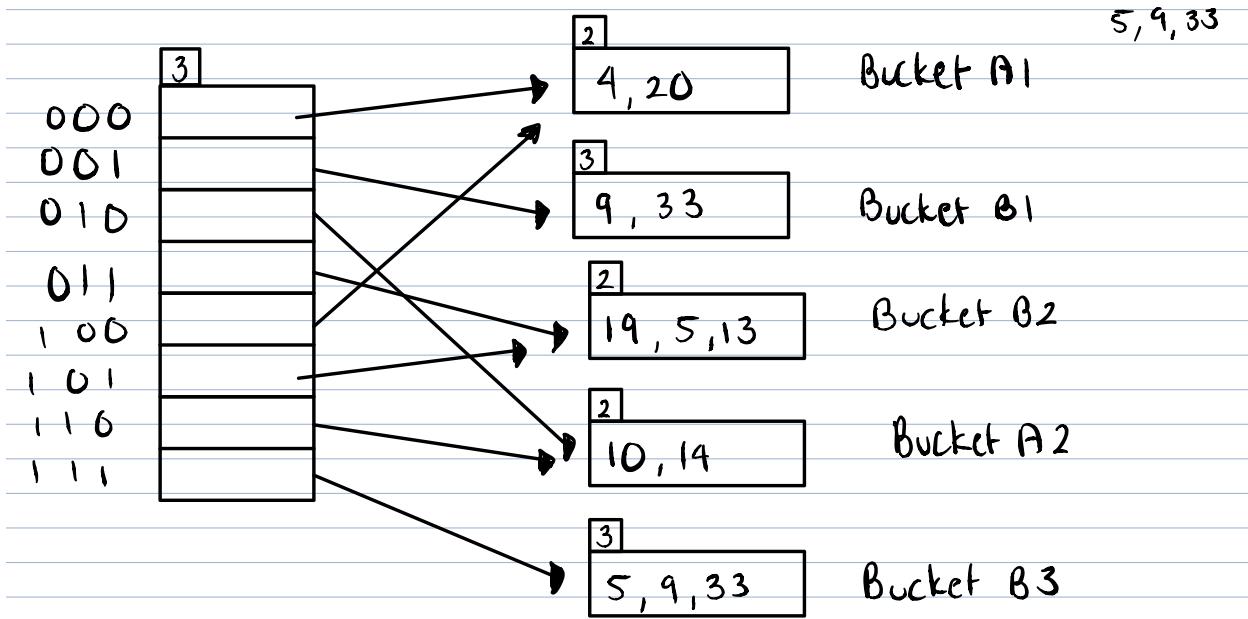
32 16 8 4 2 1

check 5 0 0 0 1 0 0, Bucket A-split



Binary conversions

	32	16	8	4	2	1	
check 4	0	0	0	1	0	0	, Bucket A1
check 10	0	0	1	0	1	0	, Bucket A2
check 20	0	1	0	1	0	0	, Bucket A1
check 5	0	0	0	1	0	1	, Bucket B1
check 9	0	0	1	0	0	1	, Bucket B1
check 19	0	1	0	0	1	1	, Bucket B2
insert 33	1	0	0	0	0	1	, Bucket B1
insert 13	0	0	1	1	0	1	, Bucket B1-split



Binary conversions

32 16 8 4 2 1

check 5 0 0 0 1 0 1 , Bucket B2

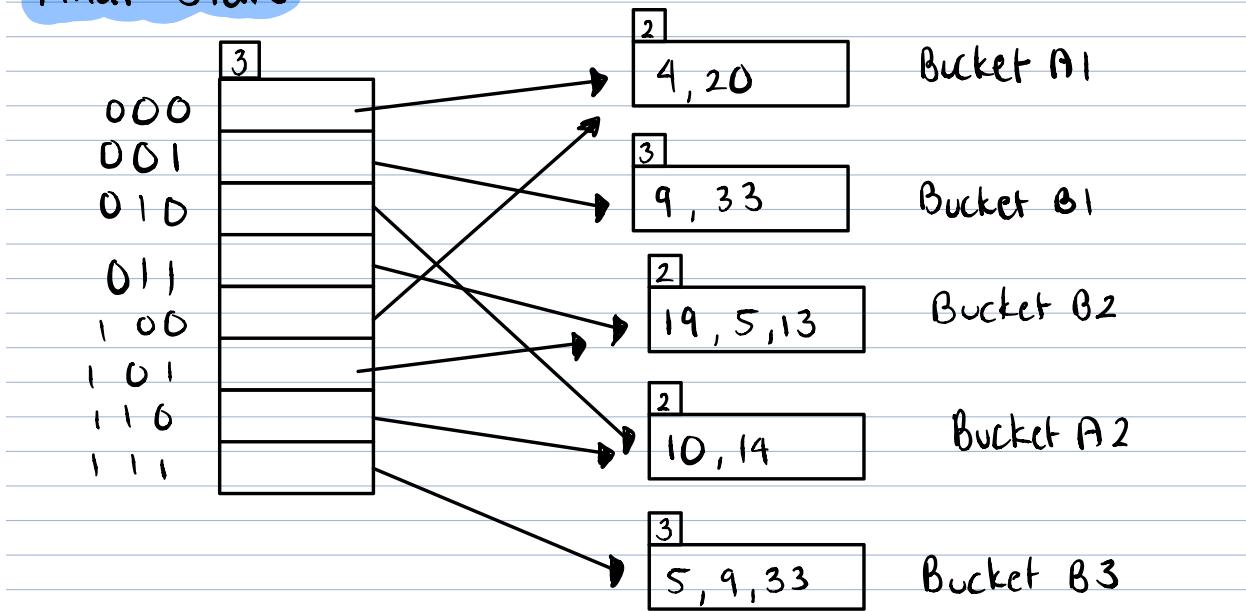
check 9 0 0 1 0 0 1 , Bucket B1

check 33 1 0 0 0 0 1 , Bucket B1

insert 13 0 0 1 1 0 1 , Bucket B2

insert 14 0 0 1 1 1 0 , Bucket A2

Final State



Problem 2

1. a. Non-blocking since it can output tuples as it processes the input.
- b. Non-blocking since column x is sorted ... once we find a varying value of x , it can output all tuples of the previous values since they make up a group.
- c. Blocking since it needs to locate elements in a group before outputting.
- d. Blocking since it needs to process all tuples before outputting the sorted list.
- e. Non-blocking since a B-tree's leaves are sorted which would mean the tuples can just be outputted as they are read.
- f. Blocking since sorting must be done, then the join.
- g. Non-blocking since we could just output the tuples as we see them.

2.

- a. Can be done in one pass with the assumption that R's distinct tuples fit in M-1 buffers, or 199 buffers.
- b. can be done in one pass following the constraint of the largest group fitting in 199 buffers.
- c. can be done in one pass as long as relation R fits in 199 buffers.
- d. cannot be done in one pass and the I/O cost would be $2 * B(R)$ for the first phase and $B(R)$ for the second phase. This would be a total of $3B(R) \rightarrow 3 * 1000 = 3000$ I/Os. Two pass external sort reads M blocks at once, then proceeds to sort, writes to disk as the first run (or phase) and finally merges the runs generating a sorted output.
- e. can be done in one pass with a total of 2,070 I/O's. It would take 2000 I/O's to read and write the sorted relation R, and 70 additional I/O's to read the index.
- f. cannot be done in one pass and in fact uses an efficient two pass algorithm with $3(B(R) + B(S))$ as an I/O cost. It would result in a total of: $3(1,000 + 150) = 3,450$ I/O's. The first phase sorts relations R and S, and the second phase merges and joins those sorted relations.
- g. can be done in one pass since relation S fits in M-1, or 199 buffers.

Problem 3

$$\begin{aligned}1. \quad Q &= T(R_1) / V(R_1, a) \\&= 400 / 50 \\&= 8\end{aligned}$$

$$\begin{aligned}2. \quad Q &= T(R_1) / V(R_1, a) * 41 \\&= 8 * 41 \\&= 328\end{aligned}$$

$$3. \quad Q = [T(R_1) / V(R_1, a) * 41] * (1/50)$$

$$= 328 * 1/50$$

$$= \boxed{7}$$

$$\begin{aligned}4. \quad Q &= T(R1) + T(R2) / \max(V(R1, b), V(R2, b)) \\&= 100 * 500 / 50 \\&= \boxed{1,000}\end{aligned}$$

$$\begin{aligned}5. \quad Q &= [T(R1) + T(R2) / \max(V(R1, b), V(R2, b))] \\&\quad * T(R3) / \max(V(R2, c), V(R3, c)) \\&= 1,000 * T(R3) / \max(V(R2, c), V(R3, c)) \\&= 1,000 * 1,000 / 100 \\&= \boxed{10,000}\end{aligned}$$

$$\begin{aligned}6. \quad Q &= [T(R1) / V(R1, a) * 41] * [T(R2) / \max(V(R1, b), V(R2, b))] \\&\quad * T(R3) / \max(V(R2, c), V(R3, c)) \\&= 328 * 500 / 50 \dots \\&= 328 * 10 \\&= 3,280 * 1,000 / 100 \\&= 3,280 * 10 \\&= \boxed{32,800}\end{aligned}$$